

Automatic Differentiation of C++ Codes With Sacado

Eric Phipps and David Gay

Sandia National Laboratories

Trilinos User's Group Meeting

November 7, 2006



Outline

- Introduction to automatic differentiation
 - Forward mode
 - Reverse mode
 - Taylor polynomial mode
- Software implementations
 - Source Transformation
 - Operator Overloading
- Sacado
 - Forward
 - Reverse
 - Higher derivatives
- Derivatives for nonlinear algorithms
- Differentiating large-scale element-based codes



What Is Automatic Differentiation (AD) ?

- All differentiable computations are composition of simple operations
 - $\sin()$, $\log()$, $+$, $*$, $/$, etc...
- We know the derivatives of these simple operations
- We have the chain rule from calculus
- Systematic application of the chain rule through your computation differentiating each statement line-by-line.



A Simple Example

$$y = \sin(e^x + x \log x), \quad x = 2$$

$$x \leftarrow 2$$

$$t_1 \leftarrow e^x$$

$$t_2 \leftarrow \log x$$

$$t_3 \leftarrow xt_2$$

$$t_4 \leftarrow t_1 + t_3$$

$$y \leftarrow \sin t_4$$

x	$\frac{d}{dx}$
2.000	1.000
7.389	7.389
0.693	0.500
1.386	1.693
8.775	9.082
0.605	-7.233

Analytic derivative evaluated to machine precision

Related Methods

$$y = \sin(e^x + x \log x), \quad x = 2$$

Automatic Differentiation

$$\begin{array}{ll}
 x \leftarrow 2 & \frac{dx}{dx} \leftarrow 1 \\
 t_1 \leftarrow e^x & \frac{dt_1}{dx} \leftarrow t_1 \frac{dx}{dx} \\
 t_2 \leftarrow \log x & \frac{dt_2}{dx} \leftarrow \frac{1}{x} \frac{dx}{dx} \\
 t_3 \leftarrow x t_2 & \frac{dt_3}{dx} \leftarrow t_2 \frac{dx}{dx} + x \frac{dt_2}{dx} \\
 t_4 \leftarrow t_1 + t_3 & \frac{dt_4}{dx} \leftarrow \frac{dt_1}{dx} + \frac{dt_3}{dx} \\
 y \leftarrow \sin t_4 & \frac{dy}{dx} \leftarrow \cos(t_4) \frac{dt_4}{dx}
 \end{array}$$

$$\frac{dy}{dx} = 7.233\,340\,400\,802\,3158$$

Symbolic Differentiation

$$\begin{aligned}
 \frac{dy}{dx} &= \cos(e^x + x \log x) \cdot \\
 &\quad (e^x + \log x + 1)
 \end{aligned}$$

$$\begin{array}{l}
 x \leftarrow 2 \\
 t_1 \leftarrow e^x \\
 t_2 \leftarrow \log x \\
 t_3 \leftarrow x t_2 \\
 t_4 \leftarrow t_1 + t_3 \\
 y \leftarrow \sin t_4
 \end{array}$$

$$\begin{array}{l}
 s_1 \leftarrow \cos t_4 \\
 s_2 \leftarrow t_1 + t_2 \\
 s_3 \leftarrow s_2 + 1
 \end{array}$$

$$\frac{dy}{dx} \leftarrow s_1 s_3$$


$$\frac{dy}{dx} = 7.233\,340\,400\,802\,3167$$

Finite Differencing

$$\begin{aligned}
 \frac{dy}{dx} &\approx \frac{y(2 + \varepsilon) - y(2)}{\varepsilon} \\
 &\approx 7.233\,343\,187
 \end{aligned}$$

Tangent Propagation

$$y = f(x), f : \mathbb{R}^n \rightarrow \mathbb{R}^m$$



- Tangents

$$y(t) = f(x(t)) \implies \dot{y} \equiv \left. \frac{dy}{dt} \right|_{t=t_0} = \frac{\partial f}{\partial x} \dot{x}$$

- For each intermediate operation

$$c = \varphi(a, b) \implies \dot{c} = \frac{\partial \varphi}{\partial a} \dot{a} + \frac{\partial \varphi}{\partial b} \dot{b}$$

- Tangents map forward through evaluation

Operation	Tangent Rule
$c = a + b$	$\dot{c} = \dot{a} + \dot{b}$
$c = a - b$	$\dot{c} = \dot{a} - \dot{b}$
$c = ab$	$\dot{c} = a\dot{b} + \dot{a}b$
$c = a/b$	$\dot{c} = (\dot{a} - c\dot{b})/b$
$c = a^b$	$\dot{c} = c(\dot{b} \log(a) + \dot{a}b/a)$
$c = \sin(a)$	$\dot{c} = \cos(a)\dot{a}$
$c = \log(a)$	$\dot{c} = \dot{a}/a$



A Simple Tangent Example

$$y_1 = \sin(e^{x_1} + x_1 x_2)$$

$$y_2 = \frac{y_1}{y_1 + x_1^2}$$

$$\begin{bmatrix} \dot{y}_1 \\ \dot{y}_2 \end{bmatrix} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} \end{bmatrix} \begin{bmatrix} \dot{x}_1 \\ \dot{x}_2 \end{bmatrix}$$

Given $x_1, x_2, \dot{x}_1, \dot{x}_2$:

$$s_1 \leftarrow e^{x_1}$$

$$\dot{s}_1 \leftarrow s_1 \dot{x}_1$$

$$s_2 \leftarrow x_1 x_2$$

$$\dot{s}_2 \leftarrow x_1 \dot{x}_2 + \dot{x}_1 x_2$$

$$s_3 \leftarrow s_1 + s_2$$

$$\dot{s}_3 \leftarrow \dot{s}_1 + \dot{s}_2$$

$$y_1 \leftarrow \sin(s_3)$$

$$\dot{y}_1 \leftarrow \cos(s_3) \dot{s}_3$$

$$s_4 \leftarrow x_1^2$$

$$\dot{s}_4 \leftarrow 2x_1 \dot{x}_1$$

$$s_5 \leftarrow y_1 + s_4$$

$$\dot{s}_5 \leftarrow \dot{y}_1 + \dot{s}_4$$

$$y_2 \leftarrow y_1 / s_5$$

$$\dot{y}_2 \leftarrow (\dot{y}_1 - y_2 \dot{s}_5) / s_5$$

Return $y_1, y_2, \dot{y}_1, \dot{y}_2$



Forward Mode AD via Tangent Propagation

- Choice of space curve $x(t)$ is arbitrary
- Tangent \dot{y} depends only on x_0, \dot{x}
- Given x_0 and v :

$$y(t) = f(x_0 + vt) \implies \dot{y} = \frac{\partial f}{\partial x_0} v \quad \text{Jacobian vector product}$$

- Propagate p vectors v_1, \dots, v_p simultaneously

$$[\dot{y}_1 \dots \dot{y}_p] = \frac{\partial f}{\partial x_0} [v_1 \dots v_p] = \frac{\partial f}{\partial x_0} V \quad \text{Jacobian matrix product}$$


- Forward mode AD:

$$(x, V) \rightarrow \left(f(x), \frac{\partial f}{\partial x} V \right)$$

- V is called the seed matrix. Setting equal to identity matrix yields full Jacobian
- Computational cost $\approx (1 + 1.5p)\text{time}(f)$

- Jacobian-vector products, directional derivatives, Jacobians for $m \geq n$

Gradient Propagation



- Gradients

$$z = \bar{y}^T y = \bar{y}^T f(x) \implies \bar{x} \equiv \left(\frac{\partial z}{\partial x} \right)^T = \left(\frac{\partial f}{\partial x} \right)^T \bar{y}$$

- For each intermediate operation

$$c = \varphi(a, b) \implies \begin{aligned} \bar{a} &= \frac{\partial z}{\partial a} = \frac{\partial z}{\partial c} \frac{\partial c}{\partial a} = \bar{c} \frac{\partial \varphi}{\partial a}, \\ \bar{b} &= \frac{\partial z}{\partial b} = \frac{\partial z}{\partial c} \frac{\partial c}{\partial b} = \bar{c} \frac{\partial \varphi}{\partial b} \end{aligned}$$

- Gradients map backward through evaluation

Operation	Gradient Rule
$c = a + b$	$\bar{a} = \bar{c}, \quad \bar{b} = \bar{c}$
$c = a - b$	$\bar{a} = \bar{c}, \quad \bar{b} = -\bar{c}$
$c = ab$	$\bar{a} = \bar{c}b, \quad \bar{b} = \bar{c}a$
$c = a/b$	$\bar{a} = \bar{c}/b, \quad \bar{b} = -\bar{c}c/b$
$c = a^b$	$\bar{a} = \bar{c}c \log(a), \quad \bar{b} = \bar{c}cb/a$
$c = \sin(a)$	$\bar{a} = \bar{c} \cos(a)$
$c = \log(a)$	$\bar{a} = \bar{c}/a$

A Simple Gradient Example

$$y_1 = \sin(e^{x_1} + x_1 x_2)$$

$$y_2 = \frac{y_1}{y_1 + x_1^2}$$

$$\begin{bmatrix} \bar{x}_1 \\ \bar{x}_2 \end{bmatrix} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} \end{bmatrix}^T \begin{bmatrix} \bar{y}_1 \\ \bar{y}_2 \end{bmatrix}$$

$$c = \varphi(a, b) \implies \begin{aligned} \bar{a} &= \bar{c} \frac{\partial \varphi}{\partial a}, \\ \bar{b} &= \bar{c} \frac{\partial \varphi}{\partial b} \end{aligned}$$

Given $x_1, x_2, \bar{y}_1, \bar{y}_2$:

$$s_1 \leftarrow e^{x_1}$$

$$s_2 \leftarrow x_1 x_2$$

$$s_3 \leftarrow s_1 + s_2$$

$$y_1 \leftarrow \sin(s_3)$$

$$s_4 \leftarrow x_1^2$$

$$s_5 \leftarrow y_1 + s_4$$

$$y_2 \leftarrow y_1 / s_5$$

$$\bar{y}_1 \leftarrow \bar{y}_1 + \bar{y}_2 / s_5, \quad \bar{s}_5 \leftarrow -y_2 \bar{y}_2 / s_5$$

$$\bar{y}_1 \leftarrow \bar{y}_1 + \bar{s}_5, \quad \bar{s}_4 \leftarrow \bar{s}_5$$

$$\bar{x}_1 \leftarrow 2\bar{s}_4 x_1$$

$$\bar{s}_3 \leftarrow \bar{y}_1 \cos(s_3)$$

$$\bar{s}_1 \leftarrow \bar{s}_3, \quad \bar{s}_2 \leftarrow \bar{s}_3$$

$$\bar{x}_1 \leftarrow \bar{x}_1 + \bar{s}_2 x_2, \quad \bar{x}_2 \leftarrow \bar{s}_2 x_1$$

$$\bar{x}_1 \leftarrow \bar{x}_1 + \bar{s}_1 s_1$$

Return $y_1, y_2, \bar{x}_1, \bar{x}_2$



Reverse Mode AD via Gradient Propagation

- Choice of normal \bar{y} is arbitrary
- Gradient \bar{x} depends only on x_0, \bar{y}
- Given x_0 and w :

$$\bar{y} = w, y = f(x) \implies \bar{x} = \left(\frac{\partial f}{\partial x} \right)^T w \quad \text{Jacobian-transpose vector product}$$

- Propagate p vectors w_1, \dots, w_p simultaneously

$$[\bar{x}_1 \dots \bar{x}_p] = \left(\frac{\partial f}{\partial x} \right)^T [w_1 \dots w_p] = \left(\frac{\partial f}{\partial x} \right)^T W \quad \text{Jacobian-transpose matrix product}$$

- Reverse mode AD:

$$(x, W) \rightarrow \left(f(x), \left(\frac{\partial f}{\partial x} \right)^T W \right)$$

- W is called the seed matrix. Setting equal to identity matrix yields full Jacobian
- Computational cost $\approx (1.5 + 2.5p)\text{time}(f)$ $m = p = 1 \implies \text{cost} \approx 4 \text{time}(f)$

- Jacobian-transpose products, gradients, Jacobians for $n > m$



Taylor Polynomial Propagation

$$y = f(x), f : \mathbb{R}^n \rightarrow \mathbb{R}^m$$

- Extension of tangent propagation to higher degree

- Given $d + 1$ coefficients $x_0, \dots, x_d \in \mathbb{R}^n$

$$x(t) \equiv \sum_{k=0}^d x_k t^k$$

$$y(t) \equiv f(x(t)) = \sum_{k=0}^d y_k t^k + O(t^{d+1})$$

$$y_k \equiv \left. \frac{1}{k!} \frac{d^k y}{dt^k} \right|_{t=0} = y_k(x_0, \dots, x_k)$$

- Computational cost $\approx O(d^2)\text{time}(f)$

Operation	Taylor Rule
$c = a + b$	$c_k = a_k + b_k$
$c = a - b$	$c_k = a_k - b_k$
$c = ab$	$c_k = \sum_{j=0}^k a_j b_{k-j}$
$c = a/b$	$c_k = \frac{1}{b_0} \left(a_k - \sum_{j=1}^k b_j c_{k-j} \right)$
$c = \exp(a)$	$c_k = \frac{1}{k} \sum_{j=1}^k j c_{k-j} a_j$
$c = \log(a)$	$c_k = \frac{1}{k a_0} \left(k a_k - \sum_{j=1}^{k-1} j a_{k-j} c_j \right)$
$s = \sin(a)$	$s_k = \frac{1}{k} \sum_{j=1}^k j a_j c_{k-j}$
$c = \cos(a)$	$c_k = -\frac{1}{k} \sum_{j=1}^k j a_j s_{k-j}$



Software Implementations

- Source transformation
 - Preprocessor reads code to be differentiated, uses AD to generate derivative code, writes-out differentiated code in original source language which is then compiled using a standard compiler
 - Resulting derivative computation is usually very efficient
 - Works well for simple languages (FORTRAN, some C)
 - ADIFOR, ADIC out of Argonne
 - Extremely difficult for C++ (no existing tool)
- Operator overloading
 - New data types are created for forward, reverse, Taylor modes, and intrinsic operations/elementary operations are overloaded to compute derivatives as a side-effect
 - Generally easy to incorporate into C++ codes
 - Generally slower than source transformation due to function call overhead
 - Requires changing data types from floats/doubles to AD types
 - C++ templates greatly help
 - ADOL-C (slow), FAD/TFAD (fast)



ADIFOR* Example

```
C      subroutine func(x, y)
C
C      double precision x(2), y(2)
C      double precision u, v, w
C
C      u = exp(x(1))
C      v = x(1)*x(2)
C      w = u+v
C      y(1) = sin(w)
C
C      u = x(1)**2
C      v = y(1) + u
C      y(2) = y(1)/v
C
C      return
C      end
```

*ADIFOR 2.0D

www-unix.mcs.anl.gov/autodiff/ADIFOR/

```
      subroutine g_func(g_p_, x, g_x, ldg_x, y, g_y, ldg_y)
C
C      Initializations removed for clarity...
C
C      d2_v = exp(x(1))
C      d1_p = d2_v
C      do g_i_ = 1, g_p_
C         g_u(g_i_) = d1_p * g_x(g_i_, 1)
C      enddo
C      u = d2_v
C-----
C      do g_i_ = 1, g_p_
C         g_v(g_i_) = x(1) * g_x(g_i_, 2) + x(2) * g_x(g_i_, 1)
C      enddo
C      v = x(1) * x(2)
C-----
C      do g_i_ = 1, g_p_
C         g_w(g_i_) = g_v(g_i_) + g_u(g_i_)
C      enddo
C      w = u + v
C-----
C      d2_v = sin(w)
C      d1_p = cos(w)
C      do g_i_ = 1, g_p_
C         g_y(g_i_, 1) = d1_p * g_w(g_i_)
C      enddo
C      y(1) = d2_v
C
C      continues...
```



(Naive) Operator Overloading Example

```
void func(const double x[], double y[]) {
    double u, v, w;
    u = exp(x[0]);
    v = x[0]*x[1];
    w = u+v;
    y[0] = sin(w);

    u = x[0]*x[0];
    v = y[0] + u;
    y[1] = y[0]/v;
}
```

```
void func(const Tangent x[], Tangent y[]) {
    Tangent u, v, w;
    u = exp(x[0]);
    v = x[0]*x[1];
    w = u+v;
    y[0] = sin(w);

    u = x[0]*x[0];
    v = y[0] + u;
    y[1] = y[0]/v;
}
```

```
class Tangent {
public:
    static const int N = 2;
    double val;
    double dot[N];
};

Tangent operator+(const Tangent& a, const Tangent& b) {
    Tangent c;
    c.val = a.val + b.val;
    for (int i=0; i<Tangent::N; i++)
        c.dot[i] = a.dot[i] + b.dot[i];
    return c;
}

Tangent operator*(const Tangent& a, const Tangent& b) {
    Tangent c;
    c.val = a.val * b.val;
    for (int i=0; i<Tangent::N; i++)
        c.dot[i] = a.val * b.dot[i] + a.dot[i]*b.val;
    return c;
}

Tangent exp(const Tangent& a) {
    Tangent c;
    c.val = exp(a.val);
    for (int i=0; i<Tangent::N; i++)
        c.dot[i] = c.val * a.dot[i];
    return c;
}
```



Introducing Sacado

(The Package Formerly Known as ADTools)

- New Trilinos package for automatic differentiation of C++ codes
- Loosely translated as “I have derived” in Spanish
- Developers: Dave Gay, Eric Phipps (with contributions from Ross Bartlett)
- Not in Trilinos 7, will be released with Trilinos 8 (Spring ‘07)
- Forward AD
 - Based on expression template-based public domain Fad package
- Reverse AD
 - Dave Gay’s Rad package
- Univariate Taylor polynomials
- Application support utilities
 - Template metaprogramming
- Coming soon
 - Ross’ ScalarFlopCounter
 - Multi-variate Taylor polynomials
 - Stochastic polynomial chaos expansions
 - Teuchos BLAS/LAPACK wrapper specializations
- Depends on Teuchos only (and currently only through tests)



The Usual Suspects

- Configure options
 - enable-sacado — Enables Sacado at Trilinos top-level
 - enable-sacado-tests, --enable-tests — Enables unit, regression, and performance tests
 - with-cppunit-prefix=[path] — Path to CppUnit for unit tests
 - with-adolc=[path] — Enables Taylor polynomial unit tests with ADOL-C
 - enable-sacado-alltests — Enables additional tests that take a VERY LONG time to compile
 - enable-sacado-examples, --enable-examples — Enables examples
 - enabled-sacado-fem-example — Enables a 1D FEM example application (additional dependencies on Epetra, NOX, LOCA, MOOCHO, Rythmos, ...)
- Mailing lists
 - Sacado-announce@software.sandia.gov
 - Sacado-checkins@software.sandia.gov
 - Sacado-developers@software.sandia.gov
 - Sacado-regression@software.sandia.gov
 - Sacado-users@software.sandia.gov
- Bugzilla: <http://software.sandia.gov/bugzilla>
- Bonsai: <http://software.sandia.gov/bonsai/cvsqueryform.cgi>
- Web: <http://software.sandai.gov/Trilinos/packages/sacado>
- Doxygen documentation



Using Sacado

- As always: `#include "Sacado.hpp"`
- All classes are templated on the Scalar type
- Forward AD classes:
 - `Sacado::Fad::DFad<ScalarT>`: Derivative array is allocated dynamically
 - `Sacado::Fad::SFad<ScalarT>`: Derivative array is allocated statically and dimension must be known at compile time
 - `Sacado::Fad::SLFad<ScalarT>`: Like SFad except allocated length may be greater than "used" length
- Reverse mode AD classes:
 - `ADvar<ScalarT>` (`Sacado_trad.h`)
- Taylor polynomial classes:
 - `Sacado::Taylor::DTaylor<ScalarT>`



sacado/example/dfad_example.cpp

```
#include "Sacado.hpp"

// The function to differentiate
template <typename ScalarT>
ScalarT func(const ScalarT& a, const ScalarT& b, const ScalarT& c) {
    ScalarT r = c*std::log(b+1.)/std::sin(a);
    return r;
}

int main(int argc, char **argv) {
    double a = std::atan(1.0);           // pi/4
    double b = 2.0;
    double c = 3.0;
    int num_deriv = 2;                  // Number of independent variables

    // Fad objects
    Sacado::Fad::DFad<double> afad(num_deriv, 0, a); // First (0) indep. var. Derivative array is [1.0 0.0]
    Sacado::Fad::DFad<double> bfad(num_deriv, 1, b); // Second (1) indep. var. Derivative array is [0.0 1.0]
    Sacado::Fad::DFad<double> cfad(c);               // Passive variable
    Sacado::Fad::DFad<double> rfad;                  // Result

    double r = func(a, b, c);                  // Compute function
    rfad = func(afad, bfad, cfad);              // Compute function and derivative with AD

    // Extract value and derivatives
    double r_ad = rfad.val();                  // r
    double drda_ad = rfad.dx(0);              // dr/da
    double drdb_ad = rfad.dx(1);              // dr/db
```



sacado/example/sfad_example.cpp

```
#include "Sacado.hpp"

// The function to differentiate
template <typename ScalarT>
ScalarT func(const ScalarT& a, const ScalarT& b, const ScalarT& c) {
    ScalarT r = c*std::log(b+1.)/std::sin(a);
    return r;
}

int main(int argc, char **argv) {
    double a = std::atan(1.0);           // pi/4
    double b = 2.0;
    double c = 3.0;
    int num_deriv = 2;                  // Number of independent variables

    // Fad objects
    Sacado::Fad::SFad<double,2> afad(num_deriv, 0, a); // First (0) indep. var. Derivative array is [1.0 0.0]
    Sacado::Fad::SFad<double,2> bfad(num_deriv, 1, b); // Second (1) indep. var. Derivative array is [0.0 1.0]
    Sacado::Fad::SFad<double,2> cfad(c);               // Passive variable
    Sacado::Fad::SFad<double,2> rfad;                  // Result

    double r = func(a, b, c);                    // Compute function
    rfad = func(afad, bfad, cfad);                // Compute function and derivative with AD

    // Extract value and derivatives
    double r_ad = rfad.val();    // r
    double drda_ad = rfad.dx(0); // dr/da
    double drdb_ad = rfad.dx(1); // dr/db
```



sacado/example/trad_example.cpp

```
#include "Sacado.hpp"

// The function to differentiate
template <typename ScalarT>
ScalarT func(const ScalarT& a, const ScalarT& b, const ScalarT& c) {
    ScalarT r = c*std::log(b+1.)/std::sin(a);
    return r;
}

int main(int argc, char **argv) {
    double a = std::atan(1.0);           // pi/4
    double b = 2.0;
    double c = 3.0;
    int num_deriv = 2;                   // Number of independent variables

    // Rad objects
    Sacado::Rad::ADvar<double> arad = a;
    Sacado::Rad::ADvar<double> brad = b;
    Sacado::Rad::ADvar<double> crad = c; // Passive variable
    Sacado::Rad::ADvar<double> rrad;     // Result
    double r = func(a, b, c);            // Compute function
    rrad = func(arad, brad, crad);        // Compute function and derivative with AD
    Sacado::Rad::ADvar<double>::Gradcomp(); // Compute gradients

    // Extract value and derivatives
    double r_ad = rrad.val();             // r
    double drda_ad = arad.adj();          // dr/da
    double drdb_ad = brad.adj();          // dr/db
```



Computing Higher Derivatives

- AD classes are templated, so AD classes can be nested to compute higher derivatives

- Forward-forward: $y = f(x) \xrightarrow{\text{for } \frac{\partial y}{\partial x} v_1} \xrightarrow{\text{for } \frac{\partial}{\partial x} \left(\frac{\partial y}{\partial x} v_1 \right)} v_2$
- Reverse-forward: $y = f(x) \xrightarrow{\text{rev } w^T \frac{\partial y}{\partial x}} \xrightarrow{\text{for } \frac{\partial}{\partial x} \left(w^T \frac{\partial y}{\partial x} \right)} v$
- Forward-Taylor: $y_0 = f(x_0) \xrightarrow{\text{for } \frac{\partial y_0}{\partial x_0} v} \xrightarrow{\text{tay } \frac{\partial y_k}{\partial x_0}} v$
- Reverse-Taylor: $y_0 = f(x_0) \xrightarrow{\text{rev } w^T \frac{\partial y_0}{\partial x_0}} \xrightarrow{\text{tay } w^T \frac{\partial y_k}{\partial x_0}}$
- Etc...



Forward or Reverse?

- Forward

- Number of independent variables \leq number of dependent variables
- Square Jacobians for Newton's method
- Sensitivities with small numbers of parameters
- Algorithm naturally calls for Jacobian-vector/matrix products
 - (Block) Matrix-free Newton-Krylov

- Reverse

- Number of independent variables $>$ number of dependent variables + 40
- Gradients of scalar valued functions
- Sensitivities with large numbers of parameters
- Algorithm naturally calls for Jacobian-transpose-vector/matrix products
 - (Block) Matrix-free solves of transpose matrix
 - Optimization



Differentiating Element-Based Codes

- Global residual computation (ignoring boundary computations):

$$f(x) = \sum_{i=1}^N Q_i^T e_{k_i}(P_i x)$$

- Jacobian computation:

$$\frac{\partial f}{\partial x} = \sum_{i=1}^N Q_i^T J_{k_i} P_i, \quad J_{k_i} = \frac{\partial e_{k_i}}{\partial x_i}, \quad x_i = P_i x$$

- Jacobian-transpose product computation:

$$w^T \frac{\partial f}{\partial x} = \sum_{i=1}^N (Q_i w)^T J_{k_i} P_i$$

- Hybrid symbolic/AD procedure
 - Element-level derivatives computed via AD
 - Exactly the same as how you would do this “manually”
 - Avoids parallelization issues




AD for Element-Level Derivatives

- Element computations are
 - Narrow and shallow
 - Dense
- Template application's element code via C++ templates
 - App developers code and maintain single templated code base
 - Easy to add new AD types
- Ideas developed within Charon
 - Large scale finite element PDE semiconductor device simulation

Performance

Set of N hypothetical chemical species:



Steady-state mass transfer equations:

$$\mathbf{u} \cdot \nabla Y_j + \nabla^2 Y_j = \dot{\omega}_j, \quad j = 1, \dots, N-1$$

$$\sum_{j=1}^N Y_j = 1$$

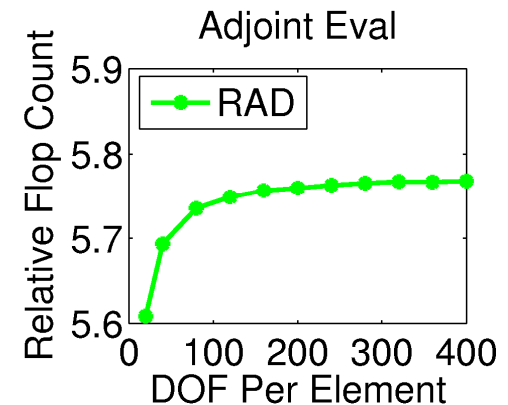


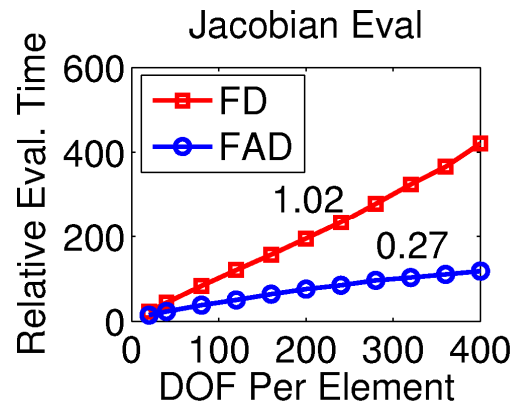
- Forward mode AD

- Faster than FD
- Better scalability in number of PDEs
- Analytic derivative
- Provides Jacobian for all Charon physics

- Reverse mode AD

- Scalable adjoint/gradient
- $$J^T w = \nabla(w^T f(x))$$

Scalability of the element-level derivative computation



DOF per element = $4 \cdot N$



How to use Sacado

- Template code to be differentiated
- Replace independent variables with AD variables
- Initialize seed matrix
- Evaluate function on AD variables
- Extract derivatives



Best Practices

- Don't differentiate your global function with AD
- Only use AD for the hard, nonlinear parts
- Do as much up-front as possible
- Never differentiate solvers with AD...instead use AD for the derivative of the solution

$$f(x, p) = 0 \implies \left(\frac{\partial f}{\partial x} \right) \frac{dx}{dp} + \frac{\partial f}{\partial p} = 0 \implies \frac{dx}{dp} = \left(\frac{\partial f}{\partial x} \right)^{-1} \frac{\partial f}{\partial p}$$

- Always put Fad inside of Rad, not the other way around



Auxiliary Slides



Why is this important?

- Speed
 - Filling a sparse Jacobian is significantly faster using forward mode AD than finite differences (FD)
 - Adjoints can be computed in a time independent of the number of independent variables using reverse mode AD
 - Higher derivatives enable more efficient algorithms
- Accuracy
 - AD produces analytic derivatives that are accurate to machine precision
 - Generally impossible to get accurate higher derivatives using FD
- Robustness
 - Analytic derivatives improve robustness of algorithms
 - Coupling of analysis algorithms requires higher derivatives
- Code maintenance
 - Hand-writing derivative code is tedious, time consuming, error prone
 - Code developers will never hand code higher derivatives
 - Using AD, code developers only need to code residuals
 - Any derivative required by an analysis tool can be computed using AD



RAD: Specialized overloading for *rf*

- Reverse mode AD tool developed by David Gay
- AD data type `ADvar`
- Forward sweep: evaluate residual fill on `ADvar` type
 - Computes values, partials
 - Store values, partials, and connectivity in “tape”
- Reverse sweep:
`ADcontext::Gradcomp()` ;
 - Accumulates adjoints
 - Reclaims “tape” memory for future gradient computations
- Block memory allocation for efficiency.
- Templated to support higher derivatives (e.g., Hessians)

$$c = \varphi(a, b) \implies \frac{\partial y_j}{\partial a} += \frac{\partial y_j}{\partial c} \frac{\partial \varphi}{\partial a}, \quad \frac{\partial y_j}{\partial b} += \frac{\partial y_j}{\partial c} \frac{\partial \varphi}{\partial b}$$

```
void ADcontext::Gradcomp() {  
    Derp *d = Derp::LastDerp;  
    d->c->aval = 1;  
    for(; d; d = d->next)  
        d->a->aval += d->da * d->c->aval;  
    // ... (arrange to recycle memory)
```

- Significantly less overhead than other approaches
 - Computes and stores partials on forward sweep
 - Only memory get/put, +, * on reverse sweep
- ADOL-C
 - Stores representation of each operation (name) and computes partial on reverse sweep
 - Allows reuse of tape



RAD Performance

Mesh quality metric from Pat Knupp

$$\tau = \det(AW^{-1})$$

$$h = 0.5(\tau + \sqrt{\tau^2 + 4\delta^2})$$

$$\mu_1 = \frac{\|AW^{-1} - I\|_F^2}{h^{2/3}}$$

A = element coordinate differences
 W = ideal element shape (fixed).

Relative Times ($f + rf$) for $f = \mu_1$	
Handcoded gradient	1.07
RAD	9.14
<i>n/c</i>	1.00
ADOL-C new tape	55.0
ADOL-C old tape	15.4

Templating Application Codes for AD

- Our interface to an application code is a templated elemental residual fill
- Templating makes it easy to interchange AD data types
- Developers only need to code residuals

```
fillResidual.h  
void fillResidual(double *x, double *f);
```

```
fillResidual.C  
void fillResidual(double *x, double *f) {  
    // Fill elemental residual  
    ...  
}
```

```
main.C  
#include "fillResidual.h"  
  
int main() {  
    double *x, *f;  
    ...  
    // Fill residual  
    fillResidual(x, f)  
    ...  
}
```



```
fillResidual.h  
template <typename ScalarT>  
void fillResidual(ScalarT *x, ScalarT *f);  
  
#include "fillResidualImpl.h"
```

```
fillResidualImpl.h  
template <typename ScalarT>  
void fillResidual(ScalarT *x, ScalarT *f) {  
    // Fill elemental residual  
    ...  
}
```

```
main.C  
#include "fillResidual.h"  
  
int main() {  
    double *x, *f;  
    Fad<double> *x_fad, *f_fad;  
    ...  
    // Fill residual  
    fillResidual<double>(x, f)  
  
    // Fill Jacobian  
    fillResidual<Fad<double>>(x_fad, f_fad);  
    ...  
}
```

Templating Application Codes for AD

- Developed reusable “template infrastructure” components:
 - Macros that explicitly instantiate template classes/functions preserving the original layout of the application code/framework into separate translation units
 - Sequence containers and iterators storing all instantiations of template classes/functions providing access to instantiations, independent of the number and types of instantiations
 - Adaptors interfacing templated AD code to non-templated derivative code (e.g., source transformation of FORTRAN)
- Eliminates a significant obstacle to widespread use of AD at Sandia
 - Easier to create/manage templated code
 - Easy to add AD types for new derivative computations

```
fillResidual.h
template <typename ScalarT>
void fillResidual(ScalarT *x, ScalarT *f);

fillResidual.C
template <typename ScalarT>
void fillResidual(ScalarT *x, ScalarT *f) {
    // Fill elemental residual
    ...
}
template void fillResidual<double>(...);
template void fillResidual< Fad<double> >(...);

main.C
#include "fillResidual.h"

int main() {
    double *x, *f;
    Fad<double> *x_fad, *f_fad;
    ...
    // Fill residual
    fillResidual<double>(x, f)
    // Fill Jacobian
    fillResidual< Fad<double> >(x_fad, f_fad);
    ...
}
```



AD With Expression Templates

Expression: $d = a + b * c$

Traditional AD code

$*$ $\left\{ \begin{array}{l} t_0 = b_0 * c_0 \\ \text{for } i = 1 : n \\ \quad t_i = b_i * c_0 + b_0 * c_i \\ \text{end} \end{array} \right.$

$+$ $\left\{ \begin{array}{l} u_0 = a_0 + t_0 \\ \text{for } i = 1 : n \\ \quad u_i = a_i + t_i \\ \text{end} \end{array} \right.$

$=$ $\left\{ \begin{array}{l} d_0 = u_0 \\ \text{for } i = 1 : n \\ \quad d_i = u_i \\ \text{end} \end{array} \right.$

Expression template code

$d_0 = a_0 + b_0 * c_0$
 $\text{for } i = 1 : n$
 $\quad d_i = a_i + b_i * c_0 + b_0 * c_i$
 end

Uses templates, lots of compiler optimization.

Requires a very good optimizing compiler!